

Solution Worksheet 2

Vishal Sontakke, Felix Späth, Michael Zellner, Christopher Zöller

1 Parallelization

We implemented the MPI.Communication via MPI.Datatypes. Our datatypes are the following ones:

```
1 MPI_Datatype _pres_right; MPI_Datatype _pres_left;
2 MPI_Datatype _pres_top; MPI_Datatype _pres_bottom;
3 MPI_Datatype _pres_front; MPI_Datatype _pres_back;
4
5 MPI_Datatype _velo_right; MPI_Datatype _velo_left;
6 MPI_Datatype _velo_top; MPI_Datatype _velo_bottom;
7 MPI_Datatype _velo_front; MPI_Datatype _velo_back;
```

The definition of the datatypes are for example the following ones (in 2D):

```
1
2 MPI_Type_vector(y,
3                 1,
4                 x+3,
5                 MY_MPI_FLOAT,
6                 &_pres_left
7                 );
8
9 MPI_Type_vector(y,
10                2,
11                x+3,
12                MY_MPI_FLOAT,
13                &_pres_right
14                );
```

By defining the datatypes like this it is not necessary to fill and read the pressure and velocity buffers. Additionally, an iterator is not necessary anymore.

1.1 Communication of Flow Quantities

- **Check out the implementation of the iterator ParallelBoundaryIterator.**
Done.
- **Read the implementation and infer its functional meaning: which cells does the iterator operate on?**
The iterator will move over the boundary layer if there is a neighboring rank available. For each cell it applies a BoundaryStencil. One can specify an offset so it can iterate over inner cells.
- **Which functions does a respective stencil object need to provide?**
The following functions are necessary because they are defined as virtual.

```
1 void applyLeftWall (FlowField & flowField, int i, int j);
2 void applyRightWall (FlowField & flowField, int i, int j);
3 void applyBottomWall (FlowField & flowField, int i, int j);
```

```

4     void applyTopWall    (FlowField & flowField, int i, int j);
5     void applyLeftWall   (FlowField & flowField, int i, int j, int k);
6     void applyRightWall  (FlowField & flowField, int i, int j, int k);
7     void applyBottomWall (FlowField & flowField, int i, int j, int k);
8     void applyTopWall    (FlowField & flowField, int i, int j, int k);
9     void applyFrontWall  (FlowField & flowField, int i, int j, int k);
10    void applyBackWall   (FlowField & flowField, int i, int j, int k);

```

- **Can you reuse existing stencil formats, or are new stencil types required?**

Since every function in BoundaryStencil is virtual we would have to define new ones.

- **Implement a boundary stencil PressureBufferFillStencil which reads the pressure values in each of the six (3D) boundary faces of a sub-domain (i.e. the domain of one process) and stores them consecutively in one-dimensional buffer arrays (one array for each of the six faces). For traversal of the respective sub-domain cells, cf. 1.**

Not necessary because we are using MPI.Datatypes.

- **Implement a boundary stencil PressureBufferReadStencil which reads data from one- dimensional arrays (one array for each of the six faces) and writes them into the correct cells of the boundary.**

Not necessary because we are using MPI.Datatypes.

- **Integrate the read- and write operations into a class PetscParallelManager. The Petsc- ParallelManager should provide a method communicatePressure() (...)**

Our communicatePressure() and communicateVelocity() methods consists of simple MPI.SendRecv operations. Filling and reading buffer is not necessary because the MPI.Datatypes will take care of that.

- **Implement the exchange of velocities analogously to the steps 1-4. Implement write- and read-stencils VelocityBufferFillStencil and VelocityBufferReadStencil and integrate them into the PetscParallelManager. The respective method communicate- Velocities() should further handle the exchange of the velocity buffers between the processes.**

Done.

- **Integrate the PetscParallelManager and respective calls to its communication methods into the class Simulation. Test the exchange of flow quantities in cavity, channel and backward-facing step flows. It is sufficient to only consider the first time step (the choice of a maximum time step should not have severe impact on the testing).**

Done. Tests were executed for 2D and 3D cases on every scenario.

1.2 Global Synchronization of the Time Step

- **Check out the implementation of setTimeStep() in the class Simulation. How is the maximum time step evaluated in the parallel simulation?**

Each rank computes its minimum timestep afterwards the global mini-

mum is derived from those values. `MPI_Allreduce(...)` takes care of the reduction and broadcast.

1.3 Validation

- **Validate your parallel implementation for different domain decompositions, that is different choices $P_x \times P_y \times P_z$, and domain sizes. Use cavity, channel and backward-facing step simulations for this purpose. Compare the solutions to your sequential program. What do you observe?**

Done.

TODO: Maybe some diagrams so that we can see the speedup

2 Scaling and Efficiency

2.1 Theory: Towards Scaling Experiments

- **How can weak and strong scaling be measured for NS-EOF? Which problems do you expect?**

Strong scaling should be easy by choosing a large scenario since NS-EOF is computebound.

Weak scaling will be tricky because the number of iterations of PETSC is dependent on the number of cells. Therefore by doubling the domainsize we expect a larger computational work.

TODO: Check again

2.2 NS-EOF in the MAC-cluster

2.2.1 Environment and Login

Done.

2.2.2 Compiling

Due to the fact that PETSC is not fully supported by the MAC cluster we compiled our own version for this worksheet. We used the version 3.7.4 the commands to compile it are the following ones:

Listing 1: Debug version

```
1 wget http://ftp.mcs.anl.gov/pub/petsc/release-snapshots/petsc-3.7.4.tar.gz
2 mkdir $HOME/petsc
3 tar -xvf petsc-3.7.4.tar.gz
4 cd petsc-3.7.4
5
6 #maccluster
7 python2 ./configure \
```

```

8  --with-blas-lapack-dir=$MKLR00T \
9  --with-mpi=1 \
10 --with-mpi-dir=$MPI_BASE \
11 --prefix=$HOME/petsc/petsc-3.7.4-debug \
12 --CC=mpiicc \
13 --CXX=mpiicpc \
14 --FC=mpiifort \
15 --CPPFLAGS="-g -O3" \ # -xHOST does not work with debugging=1
16 --CFLAGS="-g -O3" \
17 --CXXFLAGS="-g -O3" \
18 --with-vendor-compilers=intel \
19 --known-mpi-shared-libraries=1 \
20 --with-make-np=28 \
21 --with-batch=1 \
22 --with-debugging=1 \ # Set to zero for release version
23 --PETSC_ARCH=petsc-3-7-4-debug
24
25 salloc -n 1 -p snb
26 cd $HOME/petsc/petsc-3.7.4
27 ./confest-petsc-3-7-4-release
28
29 #[ Ctr + D]
30
31 ./reconfigure-petsc-3-7-4-release.py
32
33 make PETSC_DIR=$HOME/petsc-3.7.4 PETSC_ARCH=petsc-3-7-4-release all
34 make PETSC_DIR=$HOME/petsc-3.7.4 PETSC_ARCH=petsc-3-7-4-release install
35
36 #in .bashrc
37 export LD_LIBRARY_PATH=$LD_LIBRARY_PATH:$HOME/petsc/petsc-3.7.4/lib
38 export PETSC_DIR=$HOME/petsc/petsc-3.7.4

```

2.2.3 Execution

Done

2.3 Scaling Measurements

- Choose suitable simulation scenarios for the scaling measurements with respect to domain size, domain decomposition and the flow problem under consideration.

TODO:

- Run your parallel simulations for various numbers of processors and architectures. Plot your speedup and parallel efficiency in respective graphs (processes vs. speedup, e.g.). Hint: Deactivate vtk output in the scaling experiments since our output is not optimized for parallel execution yet.

TODO: Out VTK output is per rank. Moreover maybe we should compare binary and ascii output here

- How can you improve the performance of your program?

TODO:

- Optimize your software. You may work on both sequential and parallel improvements as well as on algorithmic optimizations (such as optimization of the pressure Poisson solver).

TODO:

3 Turbulence Modeling

3.1 Implementation of an algebraic turbulence model

- Compute a laminar reference case. Therefore, set $v_T = 0$.
Done with 2D.
TODO: 3D
- Implement the Prandtl mixing length model. Consider that you need to know the local distance to the nearest wall as well as S_{ij} .
Done.
- Following these instructions,
 - experiment two of the four options.
TODO:
 - explain your choices. What would you have expected and how does that differ from your findings?
TODO:
 - discuss your results and compare the results of your two modeling approaches.
TODO:
 - compare your findings to the laminar reference case.
TODO:
TODO: Ensure to base all of your discussions on appropriate data!

3.2 Implementational Details

3.2.1 Prerequisites

- Extend the configuration of the simulation such that all parameters required by your turbulence model implementation can also be read from the xml-files.
Done.

3.2.2 Turbulent viscosity and wall distance

- Extend the data structures for the turbulent case. In order to keep things simple, you may add a new scalar field for both the distance to the (nearest) wall and a scalar field to store the turbulent viscosity v_T . We further evaluate the turbulent viscosity in the cell center, i.e. at the same location as the pressure.
Done.
- Consider the basic interplay of stencils, iterators and data structures in NS-EOF. How can you apply the concept of inheritance so that you can

naturally extend your data structures and support both old and potentially new data structures in case of turbulence- model-based simulations?

Done.

- Implement a VTK visualization for the turbulent viscosity. You may base your implementations on the existing plotter (VTKStencil) for the turbulent case and create a new stencil for this purpose which operates on the new data structure.

Done.

3.2.3 Momentum equations

- Replace the discretization of the Laplacian in the evaluation of F, G and H by discrete expressions for

Done

- Improve your implementation by extending it to stretched meshes. How do you need to adapt the formulas for F, G, H from above?

TODO: Check 3D

We extend the implementation to stretched meshes by looking at the cell sizes and weighing the respective cell with the respective ratio. This is similar to what is already implemented in the code for other derivatives.

3.2.4 Parallelization

- Extend interprocess communication by communication routines for the viscosity field. The communication of turbulent viscosity is carried out analogously to the communication of the pressure. How can you use the concept of inheritance to extend and reuse the existing communication routines from the PetscParallelManager?

TODO:

- Extend the configuration by broadcasting the new parameters to all processes. You may use the methods MPI Bcast(...) and broadcast-String(...), provided by MPI and the class Configuration.

TODO:

3.2.5 Plugging things together

- Use the concept of inheritance to create a simulation class TurbulentSimulation which possesses all properties of Simulation and incorporates the features of the algebraic turbulence modeling.

Done.

- Adapt the time step restriction in TurbulentSimulation. Due to your inhomogeneous viscosity, the original expression for diffusive time step

limitations

Done.

- Implement the evaluation of the new time step criterion using the stencil-iterator concept.

Done.

4 Testing

- Test your parallel, turbulent flow simulation code in various channel flow scenarios. How can you verify that all finite difference expressions and parallel extensions are evaluated correctly?

TODO:

- Experiment with your choices on turbulence models in channel flows (cf. Sec. 3.1). What do you observe? How do the different models behave?

TODO:

- Run turbulence simulations for the backward-facing step scenario.

TODO:

TODO: Attach your discussions and appropriate plots as a pdf-file to the code you submit!